

Civic Source Scanner

System Documentation — Version 2.0

March 29, 2026

Verified Research Scaffolding for Civic Newsrooms

11 Agents • 6 Modes • 4 Kill Gates • 8 Mandatory Controls • Any US Municipality

Table of Contents

1. Executive Summary

2. Quick Start

3. Architecture & Pipeline Diagram

4. Agent Reference

5. Mandatory Controls

6. Severity Coding

7. Report Structure

8. City Configuration

9. Distribution Packager

10. Community Engagement

11. Originality Verification

12. Key Gotchas

13. Version History

14. Glossary

1. Executive Summary

Civic Source Scanner is an AI-powered research scaffolding system for civic newsrooms. It scans public records, city council agendas, meeting recordings, court filings, and community signals to produce verified story leads and reporter task memos for any US municipality.

What It Does	What It Will NOT Do
Scans 15+ official sources (Tier A/B/C) 11-agent pipeline with 4 kill gates Reporter Task Memos: confirmed facts, gaps, calls, docs, falsification Distribution packages: SEO, social, newsletters, email Persistent beat memory across runs	Publish without Tier A grounding Quote or cite Tier C sources Present AI content as photojournalism Replace reporter calls, meetings, or editorial judgment Time savings: 4-10 hrs/day. Serves: Small civic newsrooms, managing editors, investigative reporters.

2. Quick Start

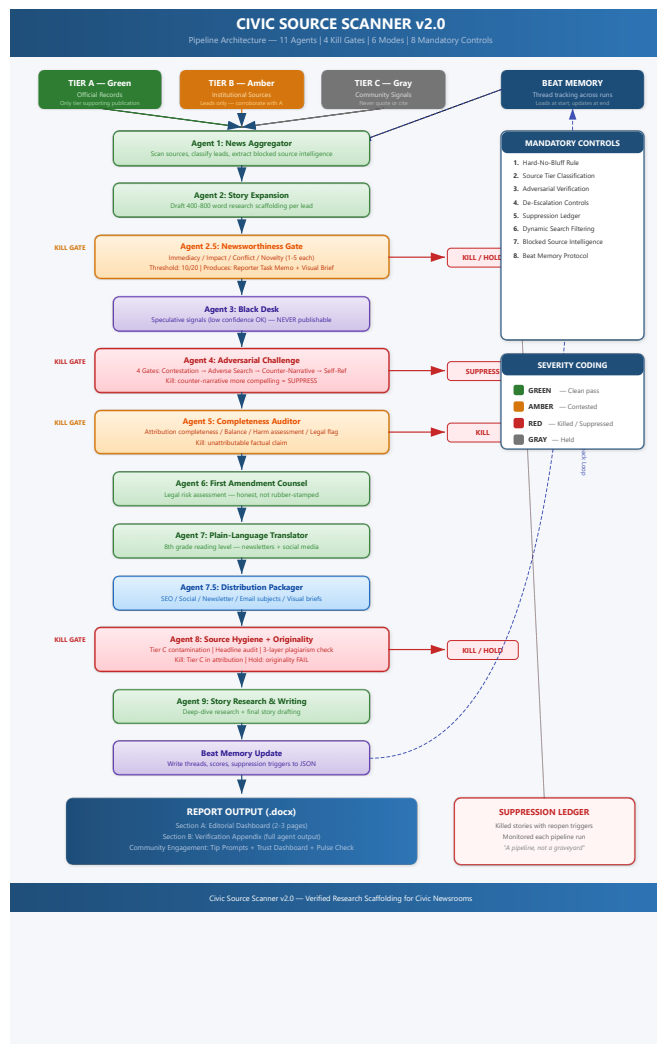
Step	Command	What Happens
1. Discover	<code>/civic-scanner discover Longmont, Colorado</code>	Builds source registry: portals, YouTube, media, forums
2. Scan	<code>/civic-scanner daily-scan Longmont</code>	Prioritized lead list with newsworthiness scores
3. Pipeline	<code>/civic-scanner full-pipeline Longmont</code>	All 11 agents, verification, task memos, distribution

All 6 Modes

Mode	Command	Output
Daily Scan	<code>daily-scan</code>	Prioritized lead list from all source tiers
Full Pipeline	<code>full-pipeline</code>	Complete story packages through all 11 agents
Verify Only	<code>verify-only {claim}</code>	4-gate adversarial verification of a specific claim
Research	<code>research {topic}</code>	Deep-dive research brief with CORA targets
Legal Threat	<code>legal-threat {situation}</code>	First Amendment counsel triage
Discover	<code>discover {city}, {state}</code>	Auto-build source registry for any US municipality

3. Architecture & Pipeline Diagram

11 sequential agents, 4 kill gates. Each agent has one editorial purpose with no overlap. Sources feed Agent 1; beat memory loads at start and updates at end; four kill gates can stop stories; surviving stories produce task memos, distribution packages, and visual briefs. Output: Section A (Editorial Dashboard) + Section B (Verification Appendix).



Source Tiers

Tier	Sources	Use
A	Gov agendas, minutes, budgets, permits, court filings, CORA, official YouTube	Only tier supporting publication
B	Newspapers, school district comms, institutional press releases	Leads only — corroborate with A
C	Reddit, Nextdoor, Facebook, YouTube comments, Twitter/X	Signals only — never quoted or cited

4. Agent Reference

#	Agent	Role	Kill Condition
1	News Aggregator	Scan sources, classify, extract 15-25 raw leads	—
2	Story Expansion	Draft 400-800 word research scaffolding per lead	—
2.5	Newsworthiness Gate	Score Immediacy/Impact/Conflict/Novelty (1-5 each); Reporter Task Memos	<10/20 = HOLD/DEMOTE
3	Black Desk	Speculative signal hunting from Tier C and blocked sources	—
4	Adversarial Challenge	4-gate: Contestation, Adverse Search, Counter-Narrative, Self-Ref	Counter-narrative wins = SUPPRESS
5	Completeness Auditor	Attribution, balance, harm, legal flag	Unattributable claim = KILL
6	1st Amendment Counsel	Legal risk assessment (honest, not rubber-stamped)	—
7	Plain-Language Translator	8th grade reading level for newsletters and social	—
7.5	Distribution Packager	SEO, social, newsletter, email subjects, visual briefs	—

8	Source Hygiene + Originality	Tier C contamination, headline audit, 3-layer plagiarism	Tier C = KILL; Orig FAIL = HOLD
9	Story Research & Writing	Deep-dive research + final story drafting	—

Kill Gates

Agent	Kill Condition	Action
2.5	Score < 10/20	HOLD (7-9) or DEMOTE (≤6)
4	Counter-narrative more compelling	SUPPRESS (red)
5	Unattributable factual claim	KILL (red)
8	Tier C in attribution / Originality FAIL	KILL (red) / HOLD

5. Mandatory Controls

1. Hard-No-Bluff Rule

No Tier A = no publication. No exceptions.

2. Source Tier Classification

A: publishable. B: leads only. C: never cite. Laundering doesn't promote.

3. Adversarial Verification

4 gates: Contestation, Adverse Search, Counter-Narrative, Self-Ref Warning.

4. De-Escalation

"Scheduled to" not "will." Fact/interpretation/allegation separated.

5. Suppression Ledger

Killed stories with reopen triggers. Pipeline, not graveyard.

6. Dynamic Search Filtering

Tier-specific allowed_domains and keep/discard criteria.

7. Blocked Source Intel

Blocked results become leads. Titles/snippets fed to Black Desk.

8. Beat Memory Protocol

Thread tracking, score trends, suppression triggers. 90-day retention.

6. Severity Coding

Color	Meaning	Example
GREEN	Clean pass, no caveats	Factual reporting of a public record, no contestation
AMBER	Passed with caveats — both sides required	Strong support but relevant counter-context
RED	Killed or suppressed	Counter-narrative won, or unattributable claim
GRAY	Held — awaiting docs or below threshold	Story scoring 8/20

7. Report Structure

Section A: Editorial Dashboard (2-3 pp)	Section B: Verification Appendix
- Pipeline stats (scanned/advanced/killed/held/suppressed) - Beat context, thread watch list, suppression triggers - Source access limitations + blocked source intel - Stories table: severity, score, beat tracking - Reporter task memos + visual direction briefs - Distribution packages - Community engagement (tips, trust dashboard, pulse)	- Full story drafts with attribution chains - Newsworthiness scoring with justifications - Black Desk speculative signals 4-gate adversarial with counter-narratives - Completeness audit + 1st Amendment analysis - Plain-language rewrites - Originality verification + source hygiene

Reporter Task Memo

Field	Content
Confirmed	Tier A grounding (specific documents)
Missing	Gaps preventing publication-ready status
Calls	Who to call (names, titles, orgs)

Documents	CORA requests or records to pull
Falsify	Evidence that would kill the story

Visual Direction Brief: Structured specs for MAP, PHOTO, CHART, INFOGRAPHIC, DIAGRAM, DOCUMENT — each with AI generation prompts for Claude Artifacts, DALL-E, matplotlib, Canva AI.

8. City Configuration

Method	Details
Auto-Discovery	<code>/civic-scanner discover Austin, Texas</code> — portals, YouTube, schools, media, forums, meeting schedules, open records
Pre-Configured	Longmont, CO (default, full registry + beat memory) • Boulder, CO
Manual	Copy <code>references/source-template.md</code> , fill sources, save as <code>references/{city}-sources.md</code>

9. Distribution Packager

Channel	Format
SEO	Headline (60 chars), meta (155 chars), keywords, URL slug
Twitter/X	280 chars + hashtags
Facebook	2-3 sentences with CTA
LinkedIn	Professional/policy framing
Nextdoor	Neighborhood-level "what this means for you"
Newsletter	50-75 word standalone + "What to watch" hook
Email	3 subject line variations (informational, curiosity, action)

10. Community Engagement

Feature	Description
Tip Intake	Per-story prompts from Task Memo gaps — tells readers exactly what the newsroom needs
Corrections	Structured log with beat memory integration; flagged in subsequent runs
Trust Dashboard	Total stories, correction rate, kill count, avg newsworthiness, Tier A %
Pulse Check	Community discussion vs pipeline coverage — prevents gov press release echo chamber

11. Originality Verification

Layer	Method	Catches
1	Structural comparison against source docs	5+ word verbatim, mirrored structure
2	Web similarity on lede paragraph	Inadvertent Tier B reproduction
3	External API (Copyscape, Originality.ai)	Broader similarity (when configured)

Verdicts: **ORIGINAL** / **FLAG** (editor review) / **FAIL** (hold for rewrite)

12. Key Gotchas

1. **Hard-No-Bluff is absolute** — no Tier A = no publish
2. **Adversarial completeness required** — evidence AGAINST is the value
3. **Self-referential warning** — AI/journalism/tech = confirmation bias risk
4. **YouTube transcripts have errors** — Tier A for content, verify quotes
5. **Astroturf detection** — manufactured pressure looks organic at Tier C
6. **SLAPP suit risk** — dev/land-use stories may trigger strategic lawsuits

- 7. **Source laundering** — C through B ≠ B
- 8. **Headline state-drift** — packaging can convert "proposed" to "approved"
- 9. **Newsworthiness ≠ sourcing** — well-sourced can be editorially thin
- 10. **Blocked ≠ absent** — blocked results contain lead intelligence
- 11. **Discovery ≠ verification** — URLs must be verified as current
- 12. **Recurring ≠ still newsworthy** — novelty decays without new facts
- 13. **Beat memory = context, not commitment** — new threads over false-matching
- 14. **Scaffolding, not journalism** — publication requires human reporting

13. Version History

Ver	Date	Feature
v1.0	2026-03-29	Initial: 5 modes, 9 agents, Hard-No-Bluff, source tiers
v1.1	2026-03-29	Dynamic Search Filtering
v1.2	2026-03-29	.docx report generation
v1.3	2026-03-29	Newsworthiness Gate, Reporter Task Memos, Headline Audit
v1.4	2026-03-29	Agent boundaries, Section A+B, severity coding
v1.5	2026-03-29	Blocked Source Intel, City Discovery Mode
v1.6	2026-03-29	Beat Memory across runs
v1.7	2026-03-29	Distribution Packager
v1.8	2026-03-29	Visual Direction Briefs + AI Generation Prompts
v1.9	2026-03-29	Community Engagement Framework
v1.10	2026-03-29	3-Layer Originality Verification
v2.0	2026-03-29	Architecture review, docs package, clean release

14. Glossary

Term	Definition
Beat Memory	Persistent JSON tracking threads, scores, suppression triggers across runs
Black Desk	Agent 3 — speculative signals; low confidence OK; never publishable
CORA	Colorado Open Records Act
Kill Gate	Pipeline checkpoint stopping stories (Agents 2.5, 4, 5, 8)
Reporter Task Memo	Confirmed / Missing / Calls / Documents / Falsify
Scaffolding	Research framework requiring human reporting before publication
Suppression Ledger	Killed stories with reopen conditions
Tier A/B/C	A: official records (publishable) / B: institutional (leads) / C: community (signals only)

Civic Source Scanner v2.0 — 30+ civic newsroom prompt variants, 4 managing editor reviews. Hard-No-Bluff enforced. Scaffolding, not journalism.