

Civic Source Scanner v2.3 — Documentation

Verified Research Scaffolding for Civic Newsrooms

March 30, 2026

1. Overview

The Civic Source Scanner is an AI-powered civic journalism pipeline built as a Claude Code skill. It scans public records, runs adversarial verification, and produces reporter-guiding intelligence for any US municipality.

- **11 agents** processing stories through scan, expansion, scoring, adversarial challenge, legal review, and distribution
- **6 modes:** daily-scan, full-pipeline, verify-only, research, legal-threat, discover
- **4 kill gates** that can stop stories before they reach an editor
- **8 mandatory controls** including the Hard-No-Bluff Rule (no Tier A source = no publication)

This system produces research scaffolding, not publishable journalism. Publication requires human reporting on top.

2. Installation

Copy skill files to your Claude Code skills directory:

```
~/ .claude/skills/civic-scanner/SKILL.md
~/ .claude/skills/civic-scanner/build-report.js
~/ .claude/skills/civic-scanner/report-schema.json
```

Install the docx npm package for report generation: `npm install docx`

3. Quick Start

- Discover sources: `/civic-scanner discover Portland, Oregon`
- Morning briefing: `/civic-scanner daily-scan --city portland`
- Full pipeline with .docx report: `/civic-scanner full-pipeline --city portland`

4. Operating Modes

Mode	Command	Description
Daily Scan	<code>daily-scan</code>	Scored morning briefing (Agents 1→2→2.5→7)
Full Pipeline	<code>full-pipeline</code>	All 11 agents with .docx report
Verify Only	<code>verify-only {claim}</code>	4-gate adversarial verification
Research	<code>research {topic}</code>	Deep-dive research brief
Legal Threat	<code>legal-threat {situation}</code>	First Amendment counsel triage
Discover	<code>discover {city}, {state}</code>	Auto-build source registry

5. Agent Pipeline

- 1. **News Aggregator** — Scan Tier A/B/C sources, classify leads
- 2. **Story Expansion** — 400-800 word drafts (full) or 200-300 word (daily)
- 3. **Newsworthiness Gate** — Score Immediacy/Impact/Conflict/Novelty; threshold 10/20; CAN KILL
- 4. **Black Desk** — Speculative signal hunting with vulnerability_type classification (7 types) and Agent 4 hand-off
- 5. **Adversarial Challenge** — 4-gate verification + Grounding Delta (auto-suppress if counter has more Tier A) + falsification search + steel-manning; CAN SUPPRESS
- 6. **Completeness Auditor** — Attribution, balance, harm assessment; CAN KILL
- 7. **First Amendment Counsel** — Legal risk + SLAPP detection + Fair Report Privilege chain verification
- 8. **Plain-Language Translator** — 8th grade reading level summaries
- 9. **Distribution Packager** — SEO, social media, newsletter, email subjects + Receipts (Tier A document IDs)
- 10. **Source Hygiene + State-Drift Linter** — Source laundering 3-step trace + hard-coded verb-tier linter + 3-layer originality check; CAN KILL/HOLD
- 11. **Story Research** — Additional deep-dive if needed

6. Kill Gates

Agent	Kill Condition	Action
Agent 2.5	Score < 10/20	HOLD (7-9) or DEMOTE (≤6)
Agent 4	Grounding Delta negative	SUPPRESS (RED) — automatic
Agent 4	Steel-man defense compelling + unaccounted for	SUPPRESS (RED)
Agent 4	Falsification search finds official denial	SUPPRESS (RED)
Agent 5	Unattributable factual claim	KILL (RED)
Agent 6	Fair Report Privilege chain broken	AMBER — mandatory rewrite
Agent 8	Tier C in attribution (including laundered)	KILL (RED)
Agent 8	Forbidden verb for source type	REJECT headline
Agent 8	Originality FAIL	HOLD for rewrite

7. Mandatory Controls

- 1. **Hard-No-Bluff Rule** — No Tier A source = no publication. Period.
- 2. **3-Tier Source Classification** — A (official records, publishable), B (leads only), C (signals, never cite)
- 3. **Adversarial Verification** — 4 gates + Grounding Delta + falsification + steel-manning
- 4. **De-Escalation** — "Scheduled to" not "will"; fact/interpretation/allegation separated
- 5. **Suppression Ledger** — Dead stories with reopen triggers; reviewed weekly
- 6. **Dynamic Search Filtering** — Tier-specific allowed_domains and keep/discard criteria
- 7. **Blocked Source Intelligence** — Blocked results become leads, not dead ends
- 8. **Beat Memory** — Persistent story thread tracking across runs; 90-day retention

8. Stress Test Suite

Test 1: Plausible Fiction (Agent 4) — Unverified memo + blog "backroom deal" claim. Expected: RED (Suppressed) via negative Grounding Delta.

Test 2: Legal/SLAPP Trap (Agent 6) — Private developer + Nextdoor "bribery" allegation. Expected: HIGH RISK flag, demand removal of characterizing language.

Test 3: State-Drift Linter (Agent 8) — Future agenda + past-tense headline "City Hikes Rates." Expected: REJECT, suggest "Council to Consider."

9. Version History

Version	Date	Feature
v1.0	2026-03-29	Initial: 5 modes, 9 agents, Hard-No-Bluff, source tiers
v1.5	2026-03-29	Blocked Source Intelligence, City Discovery Mode
v1.6	2026-03-29	Beat Memory across runs
v1.7-1.10	2026-03-29	Distribution, Visual Direction, Community Engagement, Originality
v2.0	2026-03-29	Architecture review, documentation package
v2.1	2026-03-29	Reddit tiered access, smart source lists, YouTube discovery
v2.2	2026-03-30	Enhanced daily scan (Agents 1→2→2.5→7), permanent report build script
v2.3	2026-03-30	Adversarial hardening: Grounding Delta, falsification, steel-manning, SLAPP, state-drift linter, stress tests, Receipts